

AI UT2 Question Bank

1. Define pruning? Explain alpha and beta pruning algo with an example ?

Ans: Pruning is a technique used in search algorithms, such as the minimax algorithm, to reduce the number of nodes that need to be explored. This technique involves cutting off parts of the search tree that are deemed unlikely to lead to a good solution, thereby reducing the time and computational resources required for the search.

Alpha-beta pruning is a specific form of pruning that is commonly used in game-playing algorithms. It works by maintaining two values, alpha and beta, which represent the maximum lower bound and minimum upper bound on the score that can be achieved at each node in the search tree.

During the search, alpha represents the maximum score that the maximizing player (e.g., the computer) is guaranteed to achieve at that node or above, and beta represents the minimum score that the minimizing player (e.g., the human player) is guaranteed to achieve at that node or above.

The algorithm works by recursively exploring the search tree, starting at the root node, and updating the alpha and beta values as it goes. If at any point the beta value becomes less than or equal to the alpha value, then the algorithm can stop searching that subtree, since it is impossible for the maximizing player to achieve a score higher than alpha in that subtree (or for the minimizing player to achieve a score lower than beta).

Here is an example of alpha-beta pruning in action:

Suppose we are playing a game where we can choose to roll either a 6-sided die or a 10-sided die. Our goal is to maximize the total score after rolling the dice a certain number of times. We start with a score of 0 and take turns rolling the dice.

Suppose we are at a node in the search tree where we have already rolled the 6-sided die once and have a score of 4. Our opponent (the minimizing player) can choose to roll either the 6-sided die or the 10-sided die on their turn. We consider the possible outcomes of each choice:

- If the opponent rolls the 6-sided die, they can get a score between 1 and 6, which would result in a total score for them of either 5 or 6. Since we are trying to maximize our score, we don't care about their score, so we update our alpha value to 4 (the maximum of our current score and their potential scores).

- If the opponent rolls the 10-sided die, they can get a score between 1 and 10, which would result in a total score for them of either 5 or 10. Again, we update our alpha value to 4, since we only care about our own score.

Suppose now we are at another node where we have already rolled the 6-sided die twice and have a score of 9. Our opponent can choose to roll either the 6-sided die or the 10-sided die again. We consider the possible outcomes of each choice:

- If the opponent rolls the 6-sided die, they can get a score between 1 and 6, which would result in a total score for them of either 10 or 11. However, we know from our previous search that our alpha value is already 4, so we can stop searching this subtree, since it is impossible for the opponent to achieve a score higher than 4 at this node or above.
- If the opponent rolls the 10-sided die, they can get a score between 1 and 10, which would result in a total score for them of either 19 or 10. Since our alpha value is still 4, we update our beta value to 10 (the minimum of our current score and their potential scores) and continue searching.

2. List out un-informed search algorithm and explain any one in details ?

Ans: Uninformed search algorithms, also known as blind search algorithms, are a type of search algorithm that explores the search space without any knowledge about the problem other than the problem definition. These algorithms rely on simple search strategies such as breadth-first search, depth-first search, iterative deepening search, uniform-cost search, and bidirectional search.

One of the most commonly used uninformed search algorithms is depth-first search (DFS), which is used to search a tree or graph data structure. DFS starts at the root node and explores as far as possible along each branch before backtracking.

Here is an example of how DFS works:

```
rust                                                                    Copy code

A -> B, C
B -> D, E
C -> F, G
D -> H
E -> I, J
```

If we start at node A and want to find a path to node H, we can use DFS to search the graph. Here is how the algorithm works:

1. Mark node A as visited and add it to the stack.

2. Pop node A from the stack and examine its neighbors, B and C. Mark them as visited and add them to the stack.
3. Pop node C from the stack and examine its neighbors, F and G. Mark them as visited and add them to the stack.
4. Pop node G from the stack and examine its neighbors. Since G has no unvisited neighbors, backtrack to node F.
5. Pop node F from the stack and backtrack to node C.
6. Pop node B from the stack and examine its neighbors, D and E. Mark them as visited and add them to the stack.
7. Pop node E from the stack and examine its neighbors, I and J. Mark them as visited and add them to the stack.
8. Pop node J from the stack and backtrack to node E.
9. Pop node I from the stack and backtrack to node E.
10. Pop node D from the stack and examine its neighbors, H. Since H is the goal node, return the path A -> B -> D -> H.

DFS has some advantages and disadvantages compared to other search algorithms. One advantage is that it requires less memory than breadth-first search, since it only needs to store the path from the root node to the current node, whereas breadth-first search needs to store all nodes at the current depth. However, DFS may get stuck in infinite loops or take a long time to find a solution if the search space is very large or the solution is far from the starting node.

3. List out informed search algorithm and explain A* algorithm with example ?

Ans: Informed search algorithms, also known as heuristic search algorithms, are a type of search algorithm that uses problem-specific knowledge to guide the search towards the goal state. These algorithms use a heuristic function to estimate the cost of reaching the goal from the current state. Some of the commonly used informed search algorithms are best-first search, greedy search, and A* search.

A* algorithm is a popular informed search algorithm that combines the advantages of both breadth-first search and greedy search. It uses a heuristic function to estimate the cost of reaching the goal state and combines it with the actual cost of reaching the current state to determine which node to explore next. A* algorithm always explores the node with the lowest total cost, where the total cost is the sum of the actual cost and the estimated cost.

Here's an example of how A* algorithm works:

4. Compare all algorithms according to their properties ?

Ans: There are several search algorithms that can be used to find a solution to a problem. Each algorithm has its own strengths and weaknesses. Here's a comparison of some of the commonly used search algorithms:

1. Breadth-first search (BFS): This algorithm explores all the nodes at a given depth level before moving to the next depth level. It is guaranteed to find the shortest path to the goal, but it can be slow and memory-intensive, especially for large search spaces.

2. Depth-first search (DFS): This algorithm explores as far as possible along each branch before backtracking. It is faster than BFS and requires less memory, but it may not find the shortest path to the goal.
3. Uniform-cost search (UCS): This algorithm explores the nodes in order of increasing cost, where cost is the sum of the weights of the edges along the path to the node. It is optimal and guaranteed to find the shortest path to the goal, but it can be slow and memory-intensive if the cost of the edges is high.
4. Greedy best-first search: This algorithm uses a heuristic function to estimate the cost of reaching the goal from each node and chooses the node with the lowest estimated cost. It is fast and memory-efficient, but it may not find the shortest path to the goal.
5. A* algorithm: This algorithm combines the advantages of UCS and greedy best-first search by using a heuristic function to estimate the cost of reaching the goal from each node and combining it with the actual cost of reaching the current node. It is optimal and guaranteed to find the shortest path to the goal if the heuristic function is admissible, but it can be slow and memory-intensive if the search space is very large or the heuristic function is not well-designed.
6. Depth-limited search (DLS): This algorithm is similar to DFS, but it limits the depth of the search to a certain level. It is faster and requires less memory than BFS and UCS, but it may not find the solution if the depth limit is too small.

5. how local search algorithms work explain the concept of hill climbing algorithm & simulated annealing algorithm ?

Ans: Local search algorithms are optimization algorithms that start with an initial solution and iteratively improve the solution by making small changes to it. These algorithms work by exploring the neighborhood of the current solution and selecting the best neighbor as the next solution. Local search algorithms are particularly useful when the search space is large and finding an optimal solution is computationally expensive or impossible.

Two popular local search algorithms are Hill Climbing and Simulated Annealing.

Hill Climbing Algorithm: The hill climbing algorithm is a local search algorithm that starts with an initial solution and iteratively improves the solution by making small changes to it. At each step, the algorithm evaluates the neighbors of the current solution and selects the best neighbor as the next solution. The algorithm terminates when it reaches a local optimum, which is a solution that has no better neighbors.

The hill climbing algorithm can be implemented using either the steepest ascent or the first-choice hill climbing strategy. The steepest ascent strategy evaluates all the neighbors of the current solution and selects the neighbor that has the highest objective value. The first-choice hill climbing strategy evaluates a subset of the neighbors randomly and selects the first neighbor that has a higher objective value than the current solution.

One of the main limitations of the hill climbing algorithm is that it is prone to getting stuck in local optima. This happens when the algorithm reaches a solution that has no better neighbors, but is not the global optimum.

Simulated Annealing Algorithm: The simulated annealing algorithm is a local search algorithm that uses probabilistic techniques to escape from local optima. The algorithm starts with an initial solution and iteratively improves the solution by making small changes to it. At each

step, the algorithm evaluates the neighbor of the current solution and selects it with a certain probability based on its objective value and a temperature parameter.

The temperature parameter controls the probability of accepting a worse solution than the current solution. At the beginning of the algorithm, the temperature is high, and the algorithm accepts a lot of worse solutions. As the algorithm progresses, the temperature decreases, and the algorithm becomes more selective in accepting worse solutions.

The simulated annealing algorithm terminates when the temperature reaches a minimum value or when a stopping criterion is met.

The simulated annealing algorithm can escape from local optima and find global optima by accepting worse solutions with a certain probability. However, the algorithm requires a well-designed cooling schedule to ensure that the temperature decreases at an appropriate rate.

In summary, local search algorithms such as hill climbing and simulated annealing work by iteratively improving the current solution by exploring the neighborhood of the current solution. The hill climbing algorithm selects the best neighbor as the next solution, while the simulated annealing algorithm selects a neighbor with a certain probability based on its objective value and a temperature parameter.

6. What is an adversarial search algorithm and game playing algorithm described briefly?

Ans: Adversarial search algorithms are used in game playing to find the best move in a game, where the opponent is actively trying to prevent the player from winning. These algorithms use a search strategy to explore the game tree, which represents all the possible moves that can be made by both players.

The two main types of game-playing algorithms are Minimax and Alpha-Beta pruning.

Minimax Algorithm: The minimax algorithm is a search algorithm used in game playing to find the best move in a game. The algorithm works by exploring the game tree and evaluating each possible move based on the outcome of the game. The algorithm assumes that the opponent will always play the move that maximizes their chances of winning and selects the move that minimizes the maximum loss for the player.

The minimax algorithm recursively evaluates each node in the game tree, starting from the root node, which represents the current state of the game. The algorithm computes the minimax value for each node, which is the best possible outcome for the player assuming optimal play by the opponent.

Alpha-Beta Pruning Algorithm: The alpha-beta pruning algorithm is a search algorithm used in game playing to improve the performance of the minimax algorithm. The algorithm works by pruning parts of the game tree that do not need to be evaluated, which reduces the number of nodes that need to be evaluated.

The alpha-beta pruning algorithm uses two values, alpha and beta, to determine which parts of the game tree can be pruned. Alpha represents the best score that the maximizing player has achieved so far, while beta represents the best score that the minimizing player has achieved

so far. The algorithm prunes a subtree if it can be determined that no better move exists in that subtree.

The alpha-beta pruning algorithm is more efficient than the minimax algorithm as it reduces the number of nodes that need to be evaluated. This allows the algorithm to search deeper into the game tree and find better moves.

In summary, adversarial search algorithms are used in game playing to find the best move in a game, where the opponent is actively trying to prevent the player from winning. The two main types of game-playing algorithms are Minimax and Alpha-Beta pruning. The minimax algorithm assumes that the opponent will always play the move that maximizes their chances of winning and selects the move that minimizes the maximum loss for the player. The alpha-beta pruning algorithm improves the performance of the minimax algorithm by pruning parts of the game tree that do not need to be evaluated.

7. Define following terms Unification, and resolution ?

Ans: Unification and resolution are two fundamental concepts in mathematical logic and artificial intelligence.

Unification is the process of finding a common substitution for two or more logical expressions, such that they can be made equal. The substitution is a mapping from variables to terms, where a term is a logical expression that contains no variables. For example, the substitution $\{x/a, y/b\}$ unifies the expressions $f(x, y)$ and $f(a, b)$, since substituting x with a and y with b makes the expressions equal.

Unification is a key component in many algorithms in artificial intelligence, such as theorem proving, natural language processing, and pattern matching.

Resolution is a proof procedure in mathematical logic that is used to derive new logical statements from a set of given statements. The procedure works by resolving two clauses, which are logical statements in disjunctive normal form (DNF) that are negations of each other.

The resolution rule states that if two clauses $C1$ and $C2$ contain complementary literals, then a new clause $C3$ can be derived by removing the complementary literals and joining the remaining literals with a logical OR operator. For example, if $C1$ is the clause $(\text{not } p \text{ OR } q)$ and $C2$ is the clause $(\text{not } q \text{ OR } r)$, then the clause $C3 = (\text{not } p \text{ OR } r)$ can be derived by resolving $C1$ and $C2$.

The resolution rule can be used iteratively to derive new clauses until either a contradiction is found, or no more new clauses can be derived. If a contradiction is found, then the original set of statements is inconsistent, and the proof is complete.

Resolution is an important component of automated theorem proving, which is the process of mechanically proving mathematical theorems using computers.

8. what is forward chaining and backward chaining differentiate briefly ?

Ans: Forward chaining and backward chaining are two common methods used in rule-based systems for automated reasoning and decision making.

Forward chaining, also known as data-driven reasoning or bottom-up reasoning, starts with the available data or facts and applies a set of rules to draw conclusions. The system continuously applies rules to the available data to generate new data until a desired goal or conclusion is reached. This method is often used in systems where the initial data or facts are abundant and it is desired to discover all the possible outcomes.

For example, consider a medical diagnosis system that uses forward chaining to diagnose a patient's illness. The system starts with the patient's symptoms as data and applies a set of medical rules to deduce the possible illnesses.

In contrast, backward chaining, also known as goal-driven reasoning or top-down reasoning, starts with a goal or conclusion and works backwards to find the data or facts needed to reach the goal. The system applies a set of rules to the goal to generate sub-goals, and then recursively applies rules to the sub-goals until the available data or facts are sufficient to reach the initial goal. This method is often used in systems where the desired outcome or goal is known, but the data or facts needed to reach the goal are not initially available.

For example, consider a legal decision-making system that uses backward chaining to determine the legality of a business action. The system starts with the desired outcome, which is to determine whether the business action is legal or not, and applies a set of legal rules to generate sub-goals. These sub-goals include determining the applicable laws, the facts relevant to the case, and the precedents that might be applicable.

In summary, forward chaining and backward chaining are two common methods used in rule-based systems for automated reasoning and decision making. Forward chaining starts with the available data or facts and applies a set of rules to draw conclusions, while backward chaining starts with a goal or conclusion and works backwards to find the data or facts needed to reach the goal.

9. Describe knowledge base agent ?

Ans: A knowledge-based agent is an artificial intelligence system that makes decisions by using a knowledge base, which is a set of rules and facts that the agent has about the world. The knowledge base contains information that the agent uses to reason about the current state of the world, make decisions, and take actions.

A knowledge-based agent consists of four main components:

1. Knowledge base: A set of rules and facts that the agent uses to reason about the world.
2. Inference engine: A set of algorithms and techniques that the agent uses to process the information in the knowledge base and make decisions.
3. User interface: A component that allows the agent to interact with the environment and receive input from the user.
4. Working memory: A temporary storage area where the agent stores information that it has gathered about the environment.

The knowledge base is the heart of the knowledge-based agent, as it contains the information that the agent uses to make decisions. The knowledge base can be represented in various forms, including logic-based representations, rule-based representations, and semantic networks.

The inference engine is responsible for processing the information in the knowledge base and making decisions. It uses techniques such as forward and backward chaining, constraint satisfaction, and probabilistic reasoning to make decisions.

The user interface allows the agent to interact with the environment and receive input from the user. The user interface can take various forms, including text-based interfaces, graphical user interfaces, and natural language interfaces.

The working memory is a temporary storage area where the agent stores information that it has gathered about the environment. The working memory is updated as the agent interacts with the environment, and it provides the agent with the information it needs to make decisions.

Overall, a knowledge-based agent is a powerful tool for decision making in complex environments, and it has many practical applications in areas such as finance, medicine, and law.

10. Write a short note on genetic algorithms, explain it with an example ?

Ans: Genetic algorithms (GAs) are a type of optimization algorithm inspired by the process of natural selection and genetics. These algorithms are used to solve optimization problems by mimicking the process of natural selection, crossover, and mutation.

The basic idea behind genetic algorithms is to start with a population of potential solutions, which are represented as chromosomes. Each chromosome is a vector of values that represent a potential solution to the problem. The population is then evolved over a number of generations by applying genetic operators such as selection, crossover, and mutation.

During the selection process, the best chromosomes are selected from the current population to form the next generation. This selection is usually done using fitness functions that evaluate the quality of each chromosome. The better the fitness function score of a chromosome, the higher the probability of it being selected for the next generation.

During the crossover process, pairs of chromosomes are selected and combined to produce new offspring. This process simulates the natural process of genetic recombination that occurs during sexual reproduction.

During the mutation process, small random changes are introduced into the chromosomes to simulate genetic mutations that can occur naturally. This process helps to introduce new genetic material into the population and can help to explore new regions of the search space.

As the algorithm progresses, the population evolves and the quality of the solutions improves. The algorithm terminates when a satisfactory solution is found or when a maximum number of generations is reached.

An example of a problem that can be solved using genetic algorithms is the traveling salesman problem (TSP). The TSP is a classic optimization problem that involves finding the shortest

possible route that visits a set of cities exactly once and returns to the starting city. This problem is notoriously difficult to solve, especially for large numbers of cities.

To solve the TSP using a genetic algorithm, the cities are represented as chromosomes and the fitness function evaluates the length of the route that the chromosome represents. The selection, crossover, and mutation operators are then applied to evolve the population over several generations. As the algorithm progresses, the quality of the solutions improves, and the algorithm converges to a near-optimal solution.

In summary, genetic algorithms are a type of optimization algorithm that can be used to solve a wide range of problems. These algorithms are based on the principles of natural selection and genetics, and they have been shown to be highly effective at solving complex optimization problems.

11. Explain how Genetic algorithms work. Define the terms chromosome, fitness function, crossover and mutation as used in Genetic algorithms ?

Ans: Genetic algorithms (GAs) are a type of optimization algorithm that are based on the principles of natural selection and genetics. They are used to solve optimization problems by mimicking the process of natural selection, crossover, and mutation.

The basic steps involved in a genetic algorithm are:

1. Initialization: A population of potential solutions is randomly generated. Each potential solution is represented as a chromosome, which is a vector of values that represent a potential solution to the problem.
2. Evaluation: The fitness function is applied to each chromosome in the population to determine its quality or fitness. The fitness function evaluates how well a chromosome solves the problem, and it assigns a fitness score to each chromosome.
3. Selection: The best chromosomes are selected from the current population to form the next generation. The selection is usually done using fitness functions that evaluate the quality of each chromosome. The better the fitness function score of a chromosome, the higher the probability of it being selected for the next generation.
4. Crossover: Pairs of chromosomes are selected from the current population and combined to produce new offspring. This process simulates the natural process of genetic recombination that occurs during sexual reproduction. During crossover, the genetic material of two parent chromosomes is mixed to create new offspring.
5. Mutation: Small random changes are introduced into the chromosomes to simulate genetic mutations that can occur naturally. This process helps to introduce new genetic material into the population and can help to explore new regions of the search space.
6. Termination: The algorithm terminates when a satisfactory solution is found or when a maximum number of generations is reached.

Here are some important terms used in genetic algorithms:

- Chromosome: A chromosome is a vector of values that represent a potential solution to the problem. The values in the chromosome are called genes.
- Fitness function: A fitness function is a function that evaluates how well a chromosome solves the problem. The fitness function assigns a fitness score to each chromosome based on its quality.

- Crossover: Crossover is a genetic operator that involves selecting pairs of chromosomes and combining their genetic material to produce new offspring.
- Mutation: Mutation is a genetic operator that involves making small random changes to the genes in the chromosome. This process helps to introduce new genetic material into the population and can help to explore new regions of the search space.

In summary, genetic algorithms are a powerful optimization technique that can be used to solve a wide range of problems. These algorithms mimic the natural process of evolution and genetics to find the best solution to an optimization problem. The key components of a genetic algorithm are the chromosome, fitness function, crossover, and mutation.

12. List out memory bounded algorithm. Explain SAM*.algorithm ?

Ans: Memory-bounded algorithms are algorithms that are designed to work within a fixed amount of memory. These algorithms are particularly useful in situations where the amount of available memory is limited, such as on embedded systems or mobile devices. Some common memory-bounded algorithms include:

1. Depth-limited search
2. Iterative deepening search
3. Bidirectional search
4. A* search with limited memory
5. Sampling-based motion planning algorithms

One example of a memory-bounded algorithm is SAM*, which stands for Sparse A* with Memory. SAM* is a variant of the A* search algorithm that is designed to work within a fixed amount of memory. It works by maintaining a fixed-size priority queue of nodes that have been visited, and only keeping the nodes with the best heuristic estimates.

The basic steps involved in SAM* are:

1. Initialize the start node and the priority queue with the start node.
2. While the priority queue is not empty, do the following:
 - Remove the node with the lowest f-value from the priority queue.
 - If the removed node is the goal node, return the path.
 - Generate the successors of the removed node and add them to the priority queue if they have not been visited before or if their f-value is better than the previous visit.
 - If the priority queue is full, remove the node with the highest f-value.

SAM* is particularly useful in situations where the memory is limited and the search space is very large. It is able to explore the search space efficiently while using a limited amount of memory. However, because it only keeps track of a limited number of nodes, there is a risk of missing the optimal solution.

In summary, memory-bounded algorithms are a class of algorithms that are designed to work within a fixed amount of memory. SAM* is an example of a memory-bounded algorithm that is a variant of the A* search algorithm. It is useful in situations where the amount of available memory is limited, but it may miss the optimal solution due to its limited memory capacity.